




COMP702: Classical Artificial Intelligence

2: Authored Behaviour

Agents

Agents



An **agent** is anything which perceives an **environment** through **sensors**, and acts upon that environment through **actuators**.

Performance

- ▶ An “intelligent” agent moves towards some kind of **goal**
- ▶ The goal is an **environment state** (or a set of states)
- ▶ A **performance measure** evaluates a given state for how well it fits the goal

PEAS

For each example of an agent, what are the Performance measure, Environment, Actuators and Sensors?

- ▶ A Roomba
- ▶ A self-driving car
- ▶ A chatbot
- ▶ A factory robot
- ▶ An enemy in an FPS game
- ▶ A chess AI
- ▶ A human

Types of environment

- ▶ Environments come with many different properties
- ▶ These properties influence the choice of AI architecture we use to build agents

Observability

- ▶ **Fully observable:** the agent's sensors give it full information about the state of the environment
- ▶ **Partially observable:** some aspects of the environment state are not visible to the agent's sensors
- ▶ E.g. a chess game is fully observable, a poker game is partially observable

Number of agents

- ▶ **Single agent:** our agent is the only one in the environment
- ▶ **Multi-agent:** there is more than one agent
- ▶ **Cooperative:** all agents share the same performance measure
- ▶ **Competitive:** agents' performance measures are in opposition to each other (i.e. if one agent "wins", another "loses")

Determinism

- ▶ **Deterministic**: the next state of the environment is completely determined by the current state and by the agent's action
- ▶ **Stochastic**: there is some aspect of randomness in determining the next state
- ▶ E.g. chess is deterministic; any board game involving dice rolls or random card draws is stochastic

Dynamicity

- ▶ **Static:** the environment does not change while the agent is deliberating
- ▶ **Dynamic:** the environment changes constantly
- ▶ E.g. most board games are static, most (non turn-based) video games are dynamic

Discreteness

- ▶ **Discrete:** time, percepts and actions are all discrete (from a finite set of possibilities or “integer valued”)
- ▶ **Continuous:** at least one of these is not discrete (“float valued”)
- ▶ Continuous problems are hard so we sometimes **discretise** them

Known or unknown

- ▶ Are all the details of the environment **known** to the AI designer?
- ▶ For a game or simulation: probably **yes** (unless someone else made it and we don't have the source code)
- ▶ For the real world: technically **no** (but we have physics, sociology, economics etc to give us good approximations)

Agents and AI

- ▶ The ideas of agents and environments are a useful frame for designing AI
- ▶ All(?) AI problems can be expressed in terms of creating an agent that optimises some performance measure in some environment
- ▶ Agent design boils down to: given a **percept** (and possibly some **memory** of past percepts/actions), choose the best **action** to take now

Rule-based AI

Rule-based AI

- ▶ Generally **reactive** to the state of the world
- ▶ Based on **if-then** triggers, basic **calculations**, etc.
- ▶ Generally hand-coded and only modifiable by a programmer

Case study: Ghosts in Pac-Man

- ▶ Full details: <http://gameinternals.com/understanding-pac-man-ghost-behavior>
- ▶ Each ghost has 3 states
 - ▶ Chase: head for a specific position (see next slide)
 - ▶ Scatter: head for a specific corner of the level
 - ▶ Frightened: move randomly

Ghost “personalities”

- ▶ Red ghost: aim for Pac-Man
- ▶ Pink ghost: aim for 2 spaces ahead of Pac-Man
- ▶ Blue ghost: aim for position on the line between red ghost and 2 spaces ahead of Pac-Man
- ▶ Orange ghost: aim for Pac-Man until 8 spaces away, then aim for corner

Ghost movement

- ▶ No pathfinding — greedily move towards target
- ▶ Can only change direction at an intersection
- ▶ Can't reverse or stay still
- ▶ Therefore can't get stuck, despite imperfect pathfinding

Ghost behaviour

- ▶ Behaviour rules are very simple
- ▶ However, the combination of them leads to interesting gameplay and illusion of personality

Design lessons


- ▶ AI doesn't have to be complicated
- ▶ Simple AI, when interacting with a player and each other, can give engaging results
- ▶ Bugs in AI don't always matter...

Finite state machines

Finite state machines

- ▶ A **finite state machine (FSM)** consists of:
 - ▶ A set of **states**; and
 - ▶ **Transitions** between states
- ▶ At any given time, the FSM is in a **single state**
- ▶ **Inputs** or **events (percepts)** can cause the FSM to transition to a different state
- ▶ Which state the FSM is in dictates what **actions** the agent takes

State transition diagrams



- ▶ FSMs are often drawn as **state transition diagrams**
- ▶ Reminiscent of **flowcharts** and certain types of **UML diagram**

Other uses of FSMs


As well as AI behaviours, FSMs may also be used for:

- ▶ Animation
- ▶ UI menu systems
- ▶ Dialogue trees
- ▶ Token parsing
- ▶ ...

Implementing FSMs

- ▶ Implementation needs to keep track of current state, and execute some code dependent on the state (this code itself possibly changing the current state)
- ▶ Most common approach: a big **switch-case** statement, with an **enum** type for the state
- ▶ Object-oriented approach: a `State` class, which your FSM states inherit from
- ▶ Functional approach: represent state by a function delegate
- ▶ Coroutine approach: encode your FSM logic as a procedure which runs as a coroutine (requires either refactoring logic into structured loops, or using **goto...**)

Hierarchical FSMs



- ▶ An FSM with N states has potentially N^2 transitions
- ▶ Designing complex behaviour with FSMs quickly gets unwieldy
- ▶ Hierarchical FSMs allow to group states into **super-states** to simplify defining transitions

Should you use FSMs?

- ▶ FSMs are useful for designing simple AI behaviours
- ▶ Historically an important technique for game AI
- ▶ However other techniques such as behaviour trees are more flexible and better suited to designing complex behaviours

Behaviour Trees

Behaviour trees (BTs)

- ▶ A **hierarchical** model of decision making
- ▶ Allow **complex behaviours** to be built up from **simple components**
- ▶ Allow for **more complex** behaviours than FSMs
- ▶ First used in Halo 2 (2005), now used extensively
- ▶ Also used in robotics and other non-game AI applications

Using BTs

- ▶ Fairly easy to implement; plenty of resources online
- ▶ **Unreal**: an advanced BT system is built in
- ▶ **Unity**: numerous free and paid options on the Asset Store e.g. Behavior Machine, Behavior Designer, Behave, RAIN

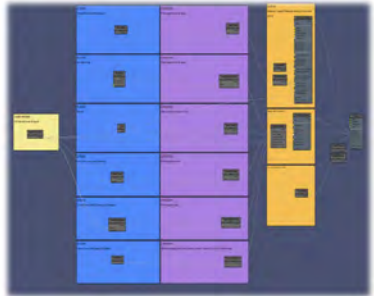
BT basics

- ▶ A BT is a **tree** of **nodes**
- ▶ On each game update (i.e. each frame), the root node is **ticked**
 - ▶ When a node is ticked, it might cause some or all of its **children** to tick as well
 - ▶ So ticks propagate down the tree from the root
- ▶ A ticked node returns one of three **statuses**:
 - ▶ Success
 - ▶ Running
 - ▶ Failure
- ▶ “Running” status allows nodes to represent operations that **last multiple frames**

Blackboard

- ▶ It is often useful to **share** data between nodes
- ▶ A **blackboard** (sometimes called a **data context**) allows this
- ▶ Blackboard defines **variables**, which can be **read** and **written** by nodes
- ▶ Blackboard can be **local** to the AI agent, **shared** between several agents, or **global** to all agents
- ▶ (Shared blackboards mean that your AI has “telepathy” — this may or may not be desirable!)

BTs in The Division



[http://www.gdcvault.com/play/1023382/
AI-Behavior-Editing-and-Debugging](http://www.gdcvault.com/play/1023382/AI-Behavior-Editing-and-Debugging)