



COMP702: Classical Artificial Intelligence

## **4: Theoretical models for games**

# **Game theory**

# Game theory

- ▶ A branch of mathematics studying **decision making**
- ▶ A **game** is a system where one or more **players** choose **actions**; the combination of these choices lead to each agent receiving a **payoff**
- ▶ Important applications in economics, ecology and social sciences as well as AI

# The Prisoner's Student's Dilemma

- ▶ Two students, **Alice** and **Bob**, are suspected of copying from each other
- ▶ Each is offered a deal in exchange for information
- ▶ Each can choose to **betray** the other or stay **silent** — but they **cannot communicate** before deciding what to do
- ▶ If **both stay silent**, both receive a C grade
- ▶ If **Alice betrays Bob**, she receives an A whilst he gets expelled
- ▶ If **Bob betrays Alice**, he receives an A whilst she gets expelled
- ▶ If **both betray each other**, both get an F

# Payoff matrix

	A silent	A betray
B silent	A: 50 B: 50	A: 70 B: -100
B betray	A: -100 B: 70	A: 0 B: 0

# Nash equilibrium

- ▶ Consider the situation where both have chosen to betray
- ▶ Neither person has anything to gain by switching to silence, assuming the other person doesn't also switch
- ▶ Such a situation is called a **Nash equilibrium**
- ▶ If all players are **rational** (in the sense of wanting to maximising payoff), they should converge upon a Nash equilibrium

# Does every game have a Nash equilibrium?

	A rock	A paper	A scissors
B rock	A: 0 B: 0	A: +1 B: -1	A: -1 B: +1
B paper	A: -1 B: +1	A: 0 B: 0	A: +1 B: -1
B scissors	A: +1 B: -1	A: -1 B: +1	A: 0 B: 0

# Nash equilibrium for Rock-Paper-Scissors

- ▶ Committing to any choice of action can be **exploited**
- ▶ E.g. if you always choose paper, I choose scissors
- ▶ If we try to reason naïvely, we get stuck in a loop
  - ▶ If I choose paper, you'll choose scissors, so I should choose rock, but then you'll choose paper, so I'll choose scissors, so you'll choose rock, so I choose paper...
- ▶ The optimum strategy is to be **unpredictable**
- ▶ Choose rock with probability  $\frac{1}{3}$ , paper with probability  $\frac{1}{3}$ , scissors with probability  $\frac{1}{3}$



# Mixed strategies

- ▶ A **mixed strategy** assigns probabilities to actions and chooses one at random
- ▶ In contrast to a **pure** or **deterministic strategy**, which always chooses the same action
- ▶ If we allow mixed strategies, **every game has at least one Nash equilibrium**

# Guess $\frac{2}{3}$ of the average

- ▶ Everyone guesses a real number (decimals are allowed) between 0 and 100 inclusive
- ▶ The winner is the person who guesses closest to  $\frac{2}{3}$  of the mean of all guesses
- ▶ Example:
  - ▶ If the guesses are 30, 40 and 80...
  - ▶ ... then the mean is  $\frac{30+40+80}{3} = 50...$
  - ▶ ... so the winning guess is 30, as this is closest to  $\frac{2}{3} \times 50 = 33.333$

# Rationality

- ▶ Rationality is a useful assumption for mathematics and AI programmers
- ▶ However it's important to remember that **humans aren't always rational**

# **Markov decision processes and games**

# Markov decision processes

- ▶ A **Markov decision process (MDP)** is defined by:
  - ▶ A finite set  $S$  of **states**;
  - ▶ A finite set  $A$  of **actions**;
  - ▶  $P(s, a, s')$  is the **probability** that action  $a$  in state  $s$  leads to state  $s'$ ;
  - ▶  $R(s, a, s')$  is the **reward** received from performing action  $a$  in state  $s$  and ending up in state  $s'$ .

# The Markov property

- ▶ Given a state  $s$  and an action  $a$ , the next state  $s'$  is determined by  $P(s, a, s')$
- ▶ The previous states before  $s$  have no effect
- ▶ Hence an MDP is “memoryless”
- ▶ (Or rather, any memory has to be contained within the state)

# (Non)determinism

- ▶ MDP defines  $P(s, a, s')$
- ▶ From state  $s$  and action  $a$ , there may be several possible states  $s'$
- ▶ MDPs are **nondeterministic** i.e. **stochastic**
- ▶ If we have

$$P(s, a, s') = \begin{cases} 1 & \text{for some } s' \\ 0 & \text{for all other } s' \end{cases}$$

then the MDP is **deterministic**

- ▶ In the deterministic case, the same state  $s$  and action  $a$  always leads to the same state  $s'$

# Markov decision problems

- ▶ A **policy** for an MDP is a function  $\pi : S \rightarrow A$
- ▶ I.e. in state  $s$ , choose action  $a = \pi(s)$
- ▶ Goal: find  $\pi$  which maximises the total reward over time




# Multi-agent MDPs

- ▶ We have assumed a **single agent** choosing a policy  $\pi$
- ▶ We can extend to  $n$  agents choosing policies  $\pi_1, \dots, \pi_n$
- ▶ The “action” is now a combination  $(a_1, \dots, a_n)$  of all the agents’ actions
- ▶ Each agent has their own reward function and is trying to maximise it
- ▶ This is a **game**!

# Minimax search

# Game trees



# Minimax

- ▶ Consider a 2-agent MDP with the following restrictions:
  - ▶ Deterministic
  - ▶ Only one agent chooses an action from a given state
  - ▶ Some states are **terminal** — these are the only ones with non-zero reward
  - ▶ The game is **zero sum**:  $R_1(s, a, s') + R_2(s, a, s') = 0$
- ▶ I want to **maximise** my reward
- ▶ My opponent wants to maximise their reward, which is the same as **minimise** my reward
- ▶ Therefore I want to **maximise** the **minimum** value my opponent can achieve

# Minimax search

- ▶ Recursively defines a **value** for non-terminal game states
- ▶ Consider each possible “next state”, i.e. each possible move
- ▶ If it's my turn, the value is the **maximum** value over next states
- ▶ If it's my opponent's turn, the value is the **minimum** value over next states

# Minimax search – example

# Minimax search pseudocode

```
procedure MINIMAX(state, currentPlayer)
  if state is terminal then
    return value of state
  else if currentPlayer = 1 then
    bestValue =  $-\infty$ 
    for each possible nextState do
       $v = \text{MINIMAX}(\text{nextState}, 3 - \text{currentPlayer})$ 
      bestValue = MAX(bestValue,  $v$ )
    return bestValue
  else if currentPlayer = 2 then
    bestValue =  $+\infty$ 
    for each possible nextState do
       $v = \text{MINIMAX}(\text{nextState}, 3 - \text{currentPlayer})$ 
      bestValue = MIN(bestValue,  $v$ )
    return bestValue
```

# Stopping early

**for each** possible nextState **do**

$v = \text{MINIMAX}(\text{nextState}, 3 - \text{currentPlayer})$

$\text{bestValue} = \text{MAX}(\text{bestValue}, v)$

- ▶ State values are always between  $-1$  and  $+1$
- ▶ So if we ever have  $\text{bestValue} = 1$ , we can stop early
- ▶ Similarly when minimising if  $\text{bestValue} = -1$
- ▶ There are techniques for smarter early stopping, e.g. alpha-beta pruning



# Using minimax search

- ▶ To decide what move to play next...
- ▶ Calculate the minimax value for each move
- ▶ Choose the move with the maximum score
- ▶ If there are several with the same score, choose one at random

# Minimax and game theory

- ▶ For a **two-player zero-sum** game with **perfect information** and **sequential moves**
- ▶ Minimax search will always find a **Nash equilibrium**
- ▶ I.e. a minimax player plays **perfectly**
- ▶ **But...**

# Minimax for larger games

- ▶ The game tree for noughts and crosses has only a few thousand states
- ▶ Most games are too large to search fully
  - ▶ Connect 4 has  $\approx 10^{13}$  states
  - ▶ Chess has  $\approx 10^{47}$  states
- ▶ Generally need to **cut off** the tree at a certain depth and use **heuristics**
- ▶ Basically, define our own reward function that (hopefully) approximates the real one

**Application to games**

# Application to board games

- ▶ The ideas of MDPs and game theory apply readily to board games
- ▶ State = state of the board
- ▶ Action = a move
- ▶ Reward = win or loss, score, etc.

# Application to video games

- ▶ *Technically* a video game is an MDP
  - ▶ State = entire state of the game (memory?)
  - ▶ Action = controller input on each frame
- ▶ However this is intractable
- ▶ Generally necessary to **abstract** parts out of the game to treat as MDPs / game theory games