



COMP702: Classical Artificial Intelligence

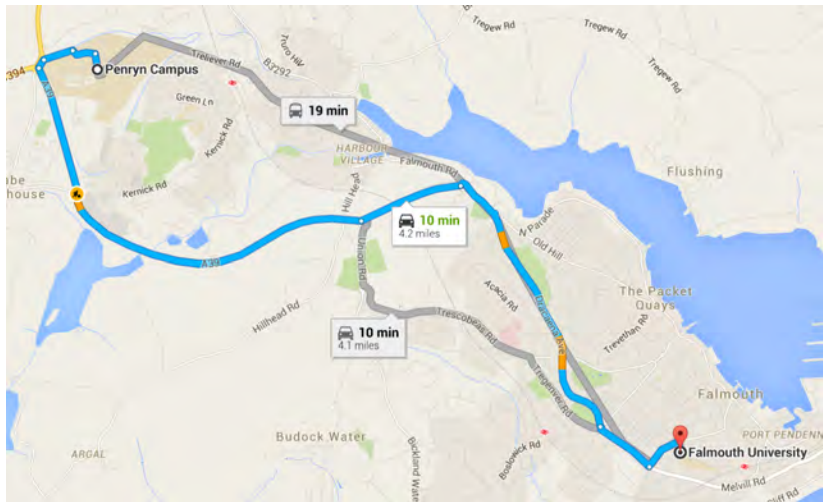
## **8: Pathfinding and Planning**

# Pathfinding


# The problem

- ▶ We have a **graph**
  - ▶ **Nodes** (points)
  - ▶ **Edges** (lines between points, each with a **length**)
- ▶ E.g. a road map
  - ▶ Nodes = addresses
  - ▶ Edges = roads
- ▶ E.g. a tile-based 2D game
  - ▶ Nodes = grid squares
  - ▶ Edges = connections between adjacent squares
- ▶ Given two nodes *A* and *B*, find the **shortest path** from *A* to *B*
  - ▶ “Shortest” in terms of edge lengths — could be distance, time, fuel cost, ...

# Applications of pathfinding



# Applications of pathfinding



# Applications of pathfinding



# Applications of pathfinding



# Applications of pathfinding





# Pathfinding as search

- ▶ Basic idea: build a **spanning tree** for the graph
- ▶ Root node is  $A$  (the start node)
- ▶ Edges in the tree are a **subset** of edges of the graph
- ▶ Once the tree includes  $B$ , we can read off the path from  $A$  to  $B$
- ▶ Need to keep track of two sets of nodes:
  - ▶ **Open set**: nodes within 1 edge of the tree, which could be added next
  - ▶ **Closed set**: nodes which have been added to the tree, and shouldn't be revisited (otherwise we could get stuck in an infinite loop)

# Graph traversal

- ▶ **Depth-first** or **breadth-first**
- ▶ Can be implemented with the open set as a **stack** or a **queue** respectively
- ▶ Inefficient — generally has to explore the **entire map**
- ▶ Finds a path, but probably not the **shortest**
- ▶ Third type of traversal: **best-first**
  - ▶ “Best” according to some heuristic evaluation
  - ▶ Often implemented with the open set as a **priority queue** — a data structure optimised for finding the **highest priority** item

# Greedy search

- ▶ Always try to move **closer** to the goal
- ▶ Visit the node whose **distance to the goal** is **minimal**
- ▶ E.g. **Euclidean** distance (straight line distance — Pythagoras' Theorem)
- ▶ Doesn't handle **dead ends** well
- ▶ Not guaranteed to find the **shortest** path

# Dijkstra's algorithm

- ▶ Let  $g(x)$  be the sum of edge weights of the path found from the start to  $x$
- ▶ Choose a node that minimises  $g(x)$
- ▶ Needs to handle cases where a shorter path to a node is discovered later in the search
- ▶ **Is** guaranteed to find the shortest path
- ▶ ... but is not the most efficient algorithm for doing so

# A\* search

- ▶ Let  $h(x)$  be an estimate of the distance from  $x$  to the goal (as in greedy search)
- ▶ Let  $g(x)$  be the distance of the path found from the start to  $x$  (as in Dijkstra's algorithm)
- ▶ Choose a node that minimises  $g(x) + h(x)$

# Properties of A\* search

- ▶ A\* is **guaranteed** to find the shortest path if the distance estimate  $h(x)$  is **admissible**
- ▶ Essentially, **admissible** means it must be an **underestimate**
  - ▶ E.g. straight line Euclidean distance is clearly an underestimate for actual travel distance
- ▶ The more accurate  $h(x)$  is, the more efficient the search
  - ▶ E.g.  $h(x) = 0$  is admissible (and gives Dijkstra's algorithm), but not very helpful
- ▶  $h(x)$  is a **heuristic**
  - ▶ In AI, a heuristic is an estimate based on human intuition
  - ▶ Heuristics are often used to prioritise search, i.e. explore the most promising options first

# Tweaking $A^*$

- ▶ Can change how  $g(x)$  is calculated
  - ▶ Increased movement cost for rough terrain, water, lava...
  - ▶ Penalty for changing direction
- ▶ Different  $h(x)$  can lead to different paths (if there are multiple “shortest” paths)

# String pulling

- ▶ Paths restricted to edges can look unnatural
- ▶ Intuition: visualise the path as a string, then pull both ends to make it taut
- ▶ Simple algorithm:
  - ▶ Found path is  $p[0], p[1], \dots, p[n]$
  - ▶ If the line from  $p[i]$  to  $p[i + 2]$  is unobstructed, remove point  $p[i + 1]$
  - ▶ Repeat until there are no more points that can be removed



# Hierarchical pathfinding in Factorio

<https://factorio.com/blog/post/fff-317>

# What is the graph?

- ▶ In a tile-based game, the graph comes from the geometry of the tiles
- ▶ In a 3D environment, the graph can be built automatically from the level geometry (e.g. with a **navigation mesh**)
- ▶ Can get complex — dynamic obstacles, vaulting, jumping, ...
- ▶ Following the path can also be complex — **steering** behaviours

**Planning**

# Planning

- ▶ An **agent** in an **environment**
- ▶ The environment has a **state**
- ▶ The agent can perform **actions** to change the state
- ▶ Actions have a **cost** associated with them
- ▶ The agent wants to change the state so as to achieve a **goal**
- ▶ Problem: find a low-cost sequence of actions that leads to the goal

# Planning as search

- ▶ We can construct a **state-action graph**
- ▶ (Similar to a **game tree**, but may include **multiple paths** or **cycles**)
- ▶ Now the planning problem becomes very similar to the pathfinding problem (albeit possibly with multiple goals)
- ▶ We can use many of the same algorithms (DFS, BFS, Dijkstra)
- ▶ We can also use  $A^*$  if we can come up with an admissible heuristic

# Representing planning problems

- ▶ We can code the state-action representation manually
- ▶ Or we can use a more general representation...

**STRIPS**

# STRIPS planning

- ▶ **S**tanford **R**esearch **I**nstitute **P**roblem **S**olver
- ▶ Describes the state of the environment by a set of **predicates** which are true
- ▶ (A predicate is basically a function which returns a `bool`)
- ▶ Models a problem as:
  - ▶ The **initial state** (a set of predicates which are true)
  - ▶ The **goal state** (a set of predicates, specifying whether each should be true or false)
  - ▶ The set of **actions**, each specifying:
    - ▶ Preconditions (a set of predicates which must be satisfied for this action to be possible)
    - ▶ Postconditions (specifying what predicates are made true or false by this action)



# STRIPS example

# STRIPS framework

- ▶ STRIPS gives a common framework for defining planning problems
- ▶ Definitions in terms of **propositional logic**
- ▶ Easy to **enumerate** and **simulate** actions, and hence search the state-action graph
- ▶ Possible to write general-purpose STRIPS solvers

**GOAP**

# GOAP

- ▶ **G**oal **O**riented **A**ction **P**lanning
- ▶ Originally developed for F.E.A.R. (2005), since used in several games
- ▶ A modified version of STRIPS specifically for real-time planning in video games

# GOAP

- ▶ Each agent has a **goal set**
  - ▶ Multiple goals with differing **priority**
  - ▶ Goals are like in STRIPS — sets of predicates that the agent wants to satisfy
- ▶ Each agent also has a set of **actions**
  - ▶ Like in STRIPS — actions have preconditions and postconditions
  - ▶ Unlike STRIPS, each action also has a **cost**

# Action sets

- ▶ Different types of agent could have the **same goals** but **different action sets**
- ▶ This will result in those agents achieving those goals in **different ways**
- ▶ NB this doesn't have to be explicitly coded — it **emerges** from the GOAP system
- ▶ E.g. this was used by the F.E.A.R. team to quickly add new enemy types

# Action sets

## Soldier

- 1 AI/Actions/Attack
- 2 AI/Actions/AttackCrouch
- 3 AI/Actions/SuppressionFire
- 4 AI/Actions/SuppressionFireFromCover
- 5 AI/Actions/FlushOutWithGrenade
- 6 AI/Actions/AttackFromCover
- 7 AI/Actions/BlindFireFromCover
- 8 AI/Actions/AttackGrenadeFromCover
- 9 AI/Actions/AttackFromView
- 10 AI/Actions/DrawWeapon
- 11 AI/Actions/HolsterWeapon
- 12 AI/Actions/ReloadCrouch
- 13 AI/Actions/ReloadCovered
- 14 AI/Actions/InspectDisturbance
- 15 AI/Actions/LookAtDisturbance
- 16 AI/Actions/SurveyArea
- 17 AI/Actions/DodgeRoll
- 18 AI/Actions/DodgeShuffle
- 19 AI/Actions/DodgeCovered
- 20 AI/Actions/Uncover
- 21 AI/Actions/AttackMelee

## Assassin


- 1 AI/Actions/Attack
- 2 AI/Actions/InspectDisturbance
- 3 AI/Actions/LookAtDisturbance
- 4 AI/Actions/SurveyArea
- 5 AI/Actions/AttackMeleeUncloaked
- 6 AI/Actions/TraverseBlockedDoor
- 7 AI/Actions/UseSmartObjectNodeMounted
- 8 AI/Actions/MountNodeUncloaked
- 9 AI/Actions/DismountNodeUncloaked
- 10 AI/Actions/TraverseLinkUncloaked
- 11 AI/Actions/AttackFromAmbush
- 12 AI/Actions/DodgeRollParanoid
- 13 AI/Actions/AttackLungeUncloaked
- 14 AI/Actions/LopeToTargetUncloaked

## Rat

- 1 AI/Actions/Animate
- 2 AI/Actions/Idle
- 3 AI/Actions/GotoNode
- 4 AI/Actions/UseSmartObjectNode

# Layering

- ▶ Goal set allows different behaviours with different priorities to be **layered**
- ▶ E.g. enemy AI in F.E.A.R.:





# Implementing GOAP

- ▶ An **abstracted** view of the game world is used for planning
- ▶ Represented as a fixed-length array (or struct) of values
- ▶ Predicates (preconditions, postconditions, goals) represented in terms of this array representation
- ▶ Most implementations also allow for **programmatic** preconditions (e.g. calling the pathfinding system to check availability of a path)

# Implementing GOAP

- ▶ Not difficult to implement
- ▶ Open-source implementations do exist
- ▶ Not built into Unity or Unreal, but asset store packages are available

# Finding the plan

- ▶ As in STRIPS, we can build a **state-action graph**
- ▶ Since actions have costs, we can use **Dijkstra's algorithm** to find the lowest cost path to the goal
- ▶ (or A\* if we can find a suitable heuristic)
- ▶ Plan is a **queue** of actions that the agent then executes
- ▶ If the plan is interrupted or fails then the agent can **replan**

# Using GOAP

- ▶ Planning is suitable when achieving a goal requires a specific **sequence of actions**
- ▶ Especially when the plan is not obvious, or when you want to let the plan be **emergent**
- ▶ Does require **abstraction** in a real-time videogame setting, though STRIPS-like definitions give a useful framework for this

# “AAA game AI” compared: GOAP vs behaviour trees

- ▶ BT: Designer specifies “how”
- ▶ GOAP: Designer specifies “what” — “how” is in whatever system is used to implement actions (FSMs in F.E.A.R.; could use BTs or hand coding)
- ▶ Both: actions (tasks in BT) are modular and reusable between agents
- ▶ GOAP: goals are also modular and reusable
- ▶ BT: goals are not represented explicitly
- ▶ BT can be classified as **authored behaviour**
- ▶ GOAP can be classified as **computational intelligence**

# Workshop